

Dimensionality Reduction

2022-06-09

R Markdown

1. Defining the question

a) Specifying the question

form the marketing department on the most relevant marketing strategies that will result in the highest no. of sales (total price including tax).

b) Defining the metrics of success

explore a recent marketing dataset by performing various unsupervised learning techniques and later providing recommendations based on your insights. reduce your dataset to a low dimensional dataset using PCA

c) Understanding the context

You are a Data analyst at Carrefour Kenya and are currently undertaking a project that will inform the marketing department on the most relevant marketing strategies that will result in the highest no. of sales (total price including tax). Your project has been divided into four parts where you'll explore a recent marketing dataset by performing various unsupervised learning techniques and later providing recommendations based on your insights.

Part 1: Dimensionality Reduction

This section of the project entails reducing your dataset to a low dimensional dataset using the t-SNE algorithm or PCA. You will be required to perform your analysis and provide insights gained from your analysis.

d) Recording the experimental design

1. Problem Definition

2. Data Sourcing
3. Check the Data
4. Perform Data Cleaning
5. Perform Exploratory Data Analysis (Univariate, Bivariate & Multivariate)
6. Implement the Solution
7. Challenge the Solution
8. Follow up Questions

e) Data relevance

The data is relevant since it was provided by the company itself and can be used to answer the question.

```
# Loading the dataset
#url <- http://bit.ly/CarreFourDataset
```

```
df <- read.csv('C:\\Users\\USER\\Documents\\Moringa School\\R Programming\\Unsupervised learning 2\\IP\\
head(df)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 1 750-67-8428      A      Member Female      Health and beauty      74.69
## 2 226-31-3081      C      Normal Female Electronic accessories      15.28
## 3 631-41-3108      A      Normal  Male      Home and lifestyle      46.33
## 4 123-19-1176      A      Member  Male      Health and beauty      58.22
## 5 373-73-7910      A      Normal  Male      Sports and travel      86.31
## 6 699-14-3026      C      Normal  Male Electronic accessories      85.39
##      Quantity      Tax      Date Time      Payment      cogs gross.margin.percentage
## 1          7 26.1415 1/5/2019 13:08      Ewallet 522.83          4.761905
## 2          5  3.8200 3/8/2019 10:29      Cash 76.40          4.761905
## 3          7 16.2155 3/3/2019 13:23 Credit card 324.31          4.761905
## 4          8 23.2880 1/27/2019 20:33      Ewallet 465.76          4.761905
## 5          7 30.2085 2/8/2019 10:37      Ewallet 604.17          4.761905
## 6          7 29.8865 3/25/2019 18:30      Ewallet 597.73          4.761905
##      gross.income Rating      Total
## 1          26.1415      9.1 548.9715
## 2           3.8200      9.6  80.2200
## 3          16.2155      7.4 340.5255
## 4          23.2880      8.4 489.0480
## 5          30.2085      5.3 634.3785
## 6          29.8865      4.1 627.6165
```

```
# Preview the last 6 items in the dataset
tail(df)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 995 652-49-6720      C      Member Female Electronic accessories      60.95
## 996 233-67-5758      C      Normal  Male      Health and beauty      40.35
## 997 303-96-2227      B      Normal Female      Home and lifestyle      97.38
## 998 727-02-1313      A      Member  Male      Food and beverages      31.84
## 999 347-56-2442      A      Normal  Male      Home and lifestyle      65.82
## 1000 849-09-3807      A      Member Female      Fashion accessories      88.34
##      Quantity      Tax      Date Time Payment      cogs gross.margin.percentage
## 995          1  3.0475 2/18/2019 11:40 Ewallet 60.95          4.761905
## 996          1  2.0175 1/29/2019 13:46 Ewallet 40.35          4.761905
## 997         10 48.6900 3/2/2019 17:16 Ewallet 973.80          4.761905
## 998          1  1.5920 2/9/2019 13:22      Cash 31.84          4.761905
## 999          1  3.2910 2/22/2019 15:33      Cash 65.82          4.761905
## 1000         7 30.9190 2/18/2019 13:28      Cash 618.38          4.761905
##      gross.income Rating      Total
## 995          3.0475      5.9  63.9975
## 996          2.0175      6.2  42.3675
## 997         48.6900      4.4 1022.4900
## 998          1.5920      7.7   33.4320
## 999          3.2910      4.1   69.1110
## 1000         30.9190      6.6  649.2990
```

```
# Checking the shape/ dimension of the dataframe
dim(df)
```

```
## [1] 1000  16
```

The dataset has 1000 rows and 16 columns

```
# Checking the summary of the dataframe
summary(df)
```

```
## Invoice.ID          Branch          Customer.type      Gender
## Length:1000        Length:1000        Length:1000        Length:1000
## Class :character    Class :character    Class :character    Class :character
## Mode :character     Mode :character     Mode :character     Mode :character
##
##
##
## Product.line        Unit.price        Quantity          Tax
## Length:1000        Min. :10.08      Min. : 1.00      Min. : 0.5085
## Class :character    1st Qu.:32.88    1st Qu.: 3.00    1st Qu.: 5.9249
## Mode :character     Median :55.23     Median : 5.00    Median :12.0880
##                     Mean :55.67      Mean : 5.51      Mean :15.3794
##                     3rd Qu.:77.94    3rd Qu.: 8.00    3rd Qu.:22.4453
##                     Max. :99.96      Max. :10.00      Max. :49.6500
##
## Date                Time                Payment          cogs
## Length:1000        Length:1000        Length:1000        Min. : 10.17
## Class :character    Class :character    Class :character    1st Qu.:118.50
## Mode :character     Mode :character     Mode :character     Median :241.76
##                     Mean :307.59
##                     3rd Qu.:448.90
##                     Max. :993.00
##
## gross.margin.percentage gross.income      Rating          Total
## Min. :4.762         Min. : 0.5085     Min. : 4.000     Min. : 10.68
## 1st Qu.:4.762       1st Qu.: 5.9249   1st Qu.: 5.500   1st Qu.: 124.42
## Median :4.762       Median :12.0880   Median : 7.000   Median : 253.85
## Mean :4.762         Mean :15.3794     Mean : 6.973     Mean : 322.97
## 3rd Qu.:4.762       3rd Qu.:22.4453   3rd Qu.: 8.500   3rd Qu.: 471.35
## Max. :4.762         Max. :49.6500     Max. :10.000     Max. :1042.65
```

```
# Checking the structure of the dataset
str(df)
```

```
## 'data.frame': 1000 obs. of 16 variables:
## $ Invoice.ID : chr "750-67-8428" "226-31-3081" "631-41-3108" "123-19-1176" ...
## $ Branch : chr "A" "C" "A" "A" ...
## $ Customer.type : chr "Member" "Normal" "Normal" "Member" ...
## $ Gender : chr "Female" "Female" "Male" "Male" ...
## $ Product.line : chr "Health and beauty" "Electronic accessories" "Home and lifestyle" ...
## $ Unit.price : num 74.7 15.3 46.3 58.2 86.3 ...
## $ Quantity : int 7 5 7 8 7 7 6 10 2 3 ...
## $ Tax : num 26.14 3.82 16.22 23.29 30.21 ...
## $ Date : chr "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" ...
## $ Time : chr "13:08" "10:29" "13:23" "20:33" ...
## $ Payment : chr "Ewallet" "Cash" "Credit card" "Ewallet" ...
## $ cogs : num 522.8 76.4 324.3 465.8 604.2 ...
## $ gross.margin.percentage: num 4.76 4.76 4.76 4.76 4.76 ...
## $ gross.income : num 26.14 3.82 16.22 23.29 30.21 ...
## $ Rating : num 9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ Total : num 549 80.2 340.5 489 634.4 ...
```

Data Cleaning

```
# Checking for uniformity in the column names of the dataset
colnames(df)
```

```
## [1] "Invoice.ID"      "Branch"
## [3] "Customer.type"   "Gender"
## [5] "Product.line"    "Unit.price"
## [7] "Quantity"        "Tax"
## [9] "Date"            "Time"
## [11] "Payment"         "cogs"
## [13] "gross.margin.percentage" "gross.income"
## [15] "Rating"          "Total"
```

```
# Checking for missing values
colSums(is.na(df))
```

```
##      Invoice.ID      Branch      Customer.type
##           0           0           0
##      Gender      Product.line      Unit.price
##           0           0           0
##      Quantity      Tax      Date
##           0           0           0
##      Time      Payment      cogs
##           0           0           0
## gross.margin.percentage      gross.income      Rating
##           0           0           0
##      Total
##           0
```

There are no missing values in the data set

```
# Checking for duplicates in the dataset
sum(duplicated(df))
```

```
## [1] 0
```

There are no duplicates in the data set

```
# Selecting the numerical values in the dataset
library(dplyr)
```


```
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##      filter, lag
## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
```

```
numeric <- select_if(df, is.numeric)
head(numeric)
```

```
##      Unit.price Quantity      Tax      cogs gross.margin.percentage gross.income
## 1      74.69         7 26.1415 522.83         4.761905         26.1415
## 2      15.28         5  3.8200  76.40         4.761905         3.8200
```

```
## 3      46.33      7 16.2155 324.31      4.761905      16.2155
## 4      58.22      8 23.2880 465.76      4.761905      23.2880
## 5      86.31      7 30.2085 604.17      4.761905      30.2085
## 6      85.39      7 29.8865 597.73      4.761905      29.8865
## Rating      Total
## 1      9.1 548.9715
## 2      9.6  80.2200
## 3      7.4 340.5255
## 4      8.4 489.0480
## 5      5.3 634.3785
## 6      4.1 627.6165
```

```
# Looking for outliers
boxplot(numeric, col = rainbow(ncol(numeric)))
```



The tax, cogs, gross income and total columns have outliers but they shall be retained for further analysis

```
# Dropping unnecessary columns
library(dplyr)
df1 = select(df, -c(Invoice.ID, Date, Time))
head(df1)
```

```
## Branch Customer.type Gender Product.line Unit.price Quantity
## 1      A      Member Female Health and beauty      74.69      7
## 2      C      Normal Female Electronic accessories      15.28      5
## 3      A      Normal  Male Home and lifestyle      46.33      7
## 4      A      Member  Male Health and beauty      58.22      8
## 5      A      Normal  Male Sports and travel      86.31      7
```

```
## 6      C      Normal   Male Electronic accessories      85.39      7
##      Tax      Payment   cogs gross.margin.percentage gross.income Rating
## 1 26.1415      Ewallet 522.83      4.761905      26.1415      9.1
## 2   3.8200      Cash   76.40      4.761905      3.8200      9.6
## 3 16.2155 Credit card 324.31      4.761905      16.2155      7.4
## 4 23.2880      Ewallet 465.76      4.761905      23.2880      8.4
## 5 30.2085      Ewallet 604.17      4.761905      30.2085      5.3
## 6 29.8865      Ewallet 597.73      4.761905      29.8865      4.1
##      Total
## 1 548.9715
## 2   80.2200
## 3 340.5255
## 4 489.0480
## 5 634.3785
## 6 627.6165
```

Exploratory Data Analysis

```
# Import necessary libraries
```

```
library(dplyr)
```

```
library(tidyr)
```

```
# Finding the mean of the numerical columns
```

```
df %>% summarise_if(is.numeric, mean)
```

```
##      Unit.price Quantity      Tax      cogs gross.margin.percentage gross.income
## 1   55.67213      5.51 15.37937 307.5874      4.761905      15.37937
##      Rating      Total
## 1 6.9727 322.9667
```

```
# Finding the median of the numerical columns
```

```
df1 %>% summarise_if(is.numeric, median)
```

```
##      Unit.price Quantity      Tax      cogs gross.margin.percentage gross.income Rating
## 1    55.23      5 12.088 241.76      4.761905      12.088      7
##      Total
## 1 253.848
```

```
# Finding the range of the numerical columns
```

```
df1 %>% summarise_if(is.numeric, range)
```

```
##      Unit.price Quantity      Tax      cogs gross.margin.percentage gross.income
## 1    10.08      1 0.5085 10.17      4.761905      0.5085
## 2    99.96     10 49.6500 993.00      4.761905      49.6500
##      Rating      Total
## 1     4    10.6785
## 2    10 1042.6500
```

```
# Finding the standard deviation of the numerical columns
```

```
df1 %>% summarise_if(is.numeric, sd)
```

```
##      Unit.price Quantity      Tax      cogs gross.margin.percentage gross.income
```

```
## 1 26.49463 2.923431 11.70883 234.1765 0 11.70883
## Rating Total
## 1 1.71858 245.8853
```

```
# Finding the variance of the numerical columns
```

```
df1 %>% summarise_if(is.numeric, var)
```

```
## Unit.price Quantity Tax cogs gross.margin.percentage gross.income
## 1 701.9653 8.546446 137.0966 54838.64 0 137.0966
## Rating Total
## 1 2.953518 60459.6
```

```
# Finding the quantiles of the numerical columns
```

```
df1 %>% summarise_if(is.numeric, quantile)
```

```
## Unit.price Quantity Tax cogs gross.margin.percentage gross.income
## 1 10.080 1 0.508500 10.1700 4.761905 0.508500
## 2 32.875 3 5.924875 118.4975 4.761905 5.924875
## 3 55.230 5 12.088000 241.7600 4.761905 12.088000
## 4 77.935 8 22.445250 448.9050 4.761905 22.445250
## 5 99.960 10 49.650000 993.0000 4.761905 49.650000
## Rating Total
## 1 4.0 10.6785
## 2 5.5 124.4224
## 3 7.0 253.8480
## 4 8.5 471.3502
## 5 10.0 1042.6500
```

```
head(df1)
```

```
## Branch Customer.type Gender Product.line Unit.price Quantity
## 1 A Member Female Health and beauty 74.69 7
## 2 C Normal Female Electronic accessories 15.28 5
## 3 A Normal Male Home and lifestyle 46.33 7
## 4 A Member Male Health and beauty 58.22 8
## 5 A Normal Male Sports and travel 86.31 7
## 6 C Normal Male Electronic accessories 85.39 7
## Tax Payment cogs gross.margin.percentage gross.income Rating
## 1 26.1415 Ewallet 522.83 4.761905 26.1415 9.1
## 2 3.8200 Cash 76.40 4.761905 3.8200 9.6
## 3 16.2155 Credit card 324.31 4.761905 16.2155 7.4
## 4 23.2880 Ewallet 465.76 4.761905 23.2880 8.4
## 5 30.2085 Ewallet 604.17 4.761905 30.2085 5.3
## 6 29.8865 Ewallet 597.73 4.761905 29.8865 4.1
## Total
## 1 548.9715
## 2 80.2200
## 3 340.5255
## 4 489.0480
## 5 634.3785
## 6 627.6165
```


```
# Countplots for the categorical variables
```


```
library(ggplot2)
cat = df1[, c(1:4,8)]
```


```

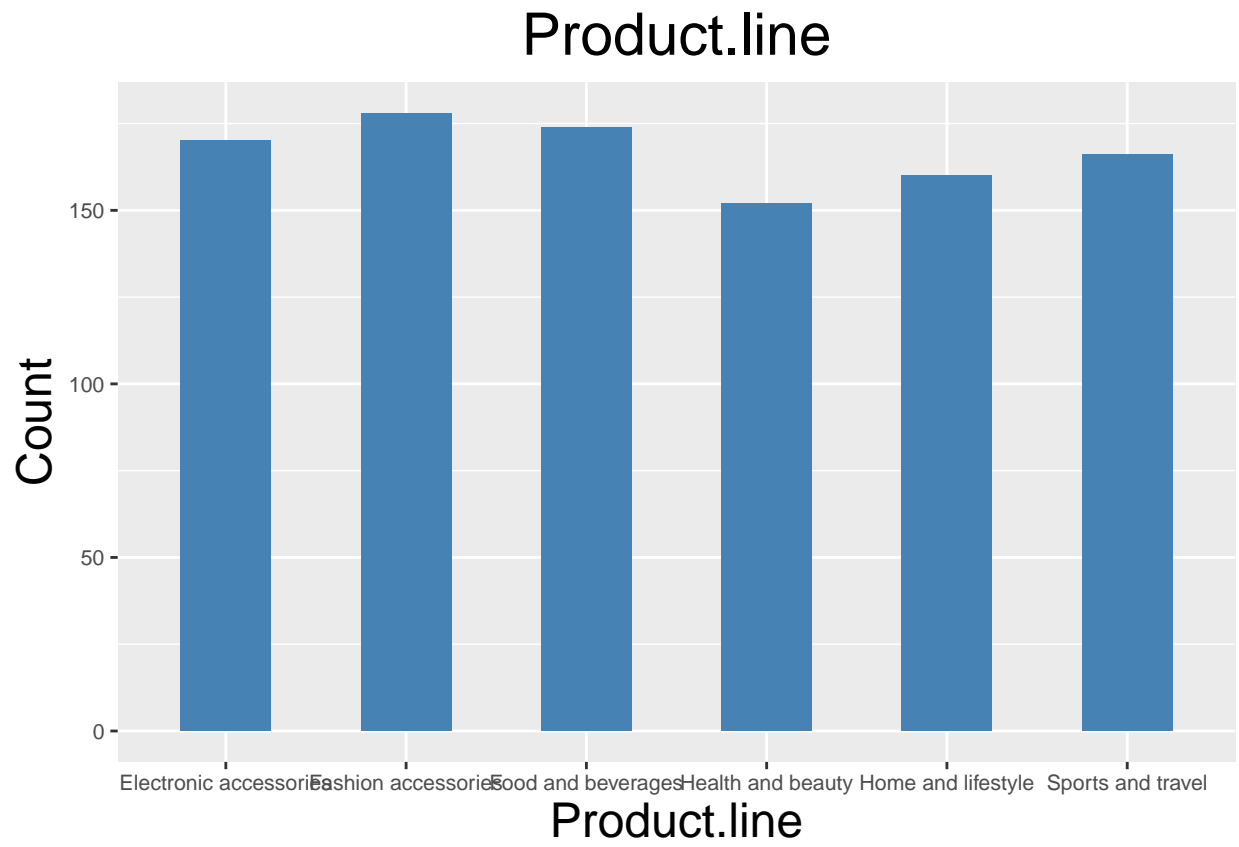
for (i in 1:length(cat)) {
  options(repr.plot.width = 30, repr.plot.height = 8)
  print(ggplot(cat, aes(x = factor(cat[,i]))) +
    geom_bar(fill = "steelblue", width = 0.5) +
    labs(title = names(cat[i]), x = names(cat[i]), y = "Count") +
    theme(axis.text = element_text(size=8),
          axis.title = element_text(size = 18),
          plot.title = element_text(hjust = 0.5, size = 22)))
  cat("\n", "\n")
}


```











```
# Checking the actual count in gender
table(df1$Gender)
```

```
##
## Female    Male
##      501    499
```

```
# Checking the actual count in gender
table(df1$Customer.type)
```

```
##
## Member Normal
##      501    499
```

Most of the payment was done via e-wallet, followed by cash then credit card. Most clients used branch A. There was a difference of 2 people between the members and normal people. There were 2 more females than males. Fashion accessories accounted for the majority of the sales, while health and beauty for the least.

Bivariate Analysis

```
# Installing GGally package to plot the pairplot
library("GGally")
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
head(df1)
```

```
##   Branch Customer.type Gender      Product.line Unit.price Quantity
## 1     A      Member Female  Health and beauty    74.69         7
## 2     C      Normal Female Electronic accessories  15.28         5
## 3     A      Normal   Male   Home and lifestyle   46.33         7
## 4     A      Member   Male   Health and beauty   58.22         8
## 5     A      Normal   Male   Sports and travel   86.31         7
## 6     C      Normal   Male Electronic accessories  85.39         7
##      Tax      Payment   cogs gross.margin.percentage gross.income Rating
## 1 26.1415      Ewallet 522.83          4.761905      26.1415      9.1
## 2  3.8200       Cash  76.40          4.761905       3.8200      9.6
## 3 16.2155 Credit card 324.31          4.761905      16.2155      7.4
## 4 23.2880      Ewallet 465.76          4.761905      23.2880      8.4
## 5 30.2085      Ewallet 604.17          4.761905      30.2085      5.3
## 6 29.8865      Ewallet 597.73          4.761905      29.8865      4.1
##      Total
## 1 548.9715
## 2  80.2200
## 3 340.5255
## 4 489.0480
## 5 634.3785
## 6 627.6165
```

```
# Plotting pair plots for numeric columns
```

```
options(repr.plot.width = 40, repr.plot.height = 18)
ggpairs(df1[, c(5:7, 9:13)], upper = list(continuous = wrap("cor", size = 7))) +
  labs(title = "Pairwise plots of numeric columns in the dataset") +
  theme_grey(base_size = 10) +
  theme(plot.title = element_text(hjust = 0.3))
```

```
## Warning in cor(x, y): the standard deviation is zero
```

```
## Warning in cor(x, y): the standard deviation is zero
```

```
## Warning in cor(x, y): the standard deviation is zero
```

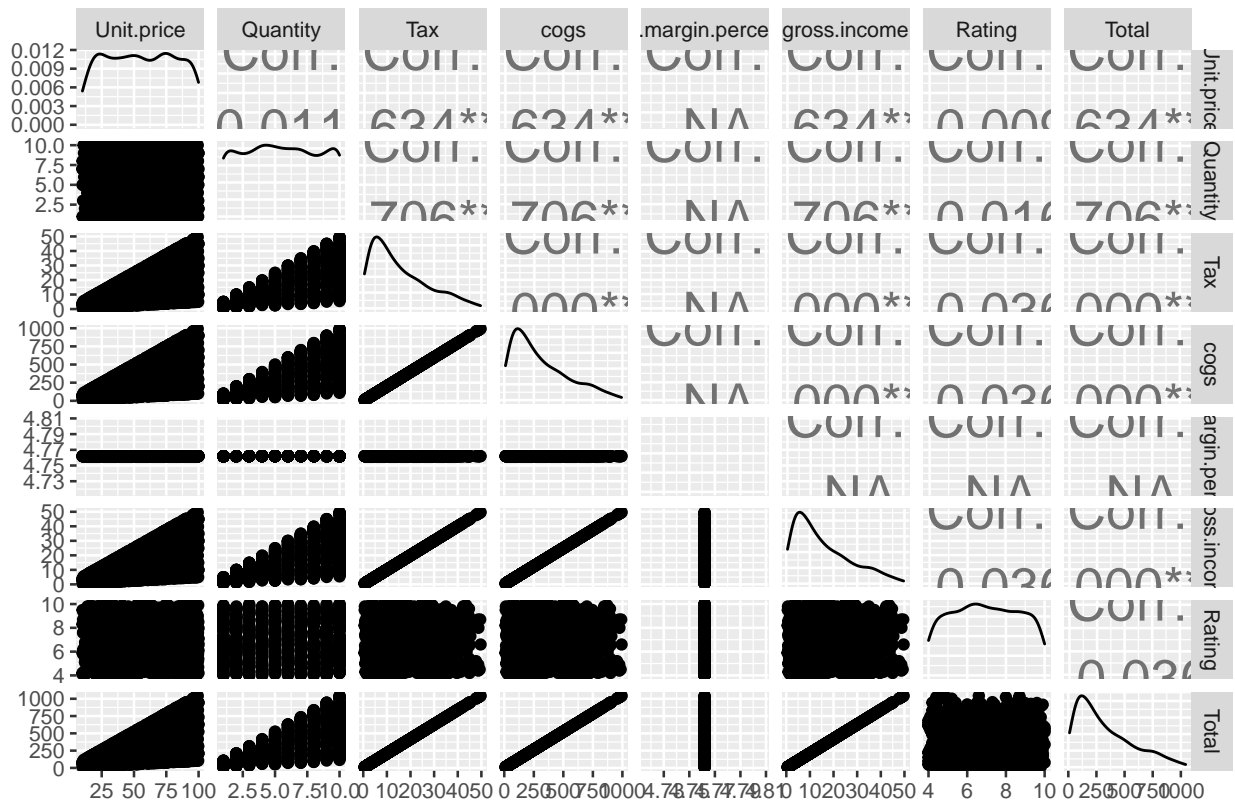
```
## Warning in cor(x, y): the standard deviation is zero
```

```
## Warning in cor(x, y): the standard deviation is zero
```

```
## Warning in cor(x, y): the standard deviation is zero
```

```
## Warning in cor(x, y): the standard deviation is zero
```

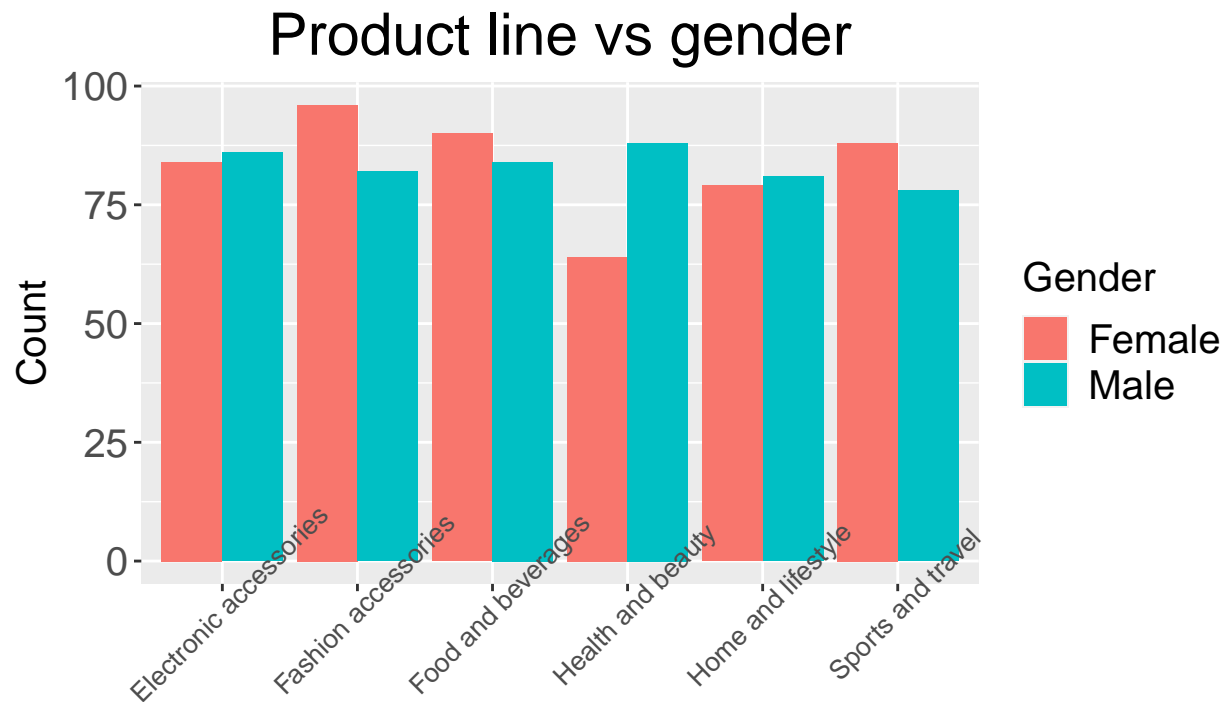
Pairwise plots of numeric columns in the dataset



There is a positive correlation between tax and cogs, tax and gross income, tax and total, cogs and total

Plot of product line vs gender

```
options(repr.plot.width = 20, repr.plot.height = 10)
ggplot(df1, aes(x = Product.line, fill = Gender)) +
  geom_bar(position = "dodge") +
  labs(title = "Product line vs gender", x = "Product line", y = "Count") +
  theme(axis.text.x = element_text(size=10, angle = 45),
        axis.text.y = element_text(size=15),
        axis.title = element_text(size = 15),
        plot.title = element_text(hjust = 0.5, size = 22),
        legend.title = element_text(size=15),
        legend.text = element_text(size=15))
```



Product line

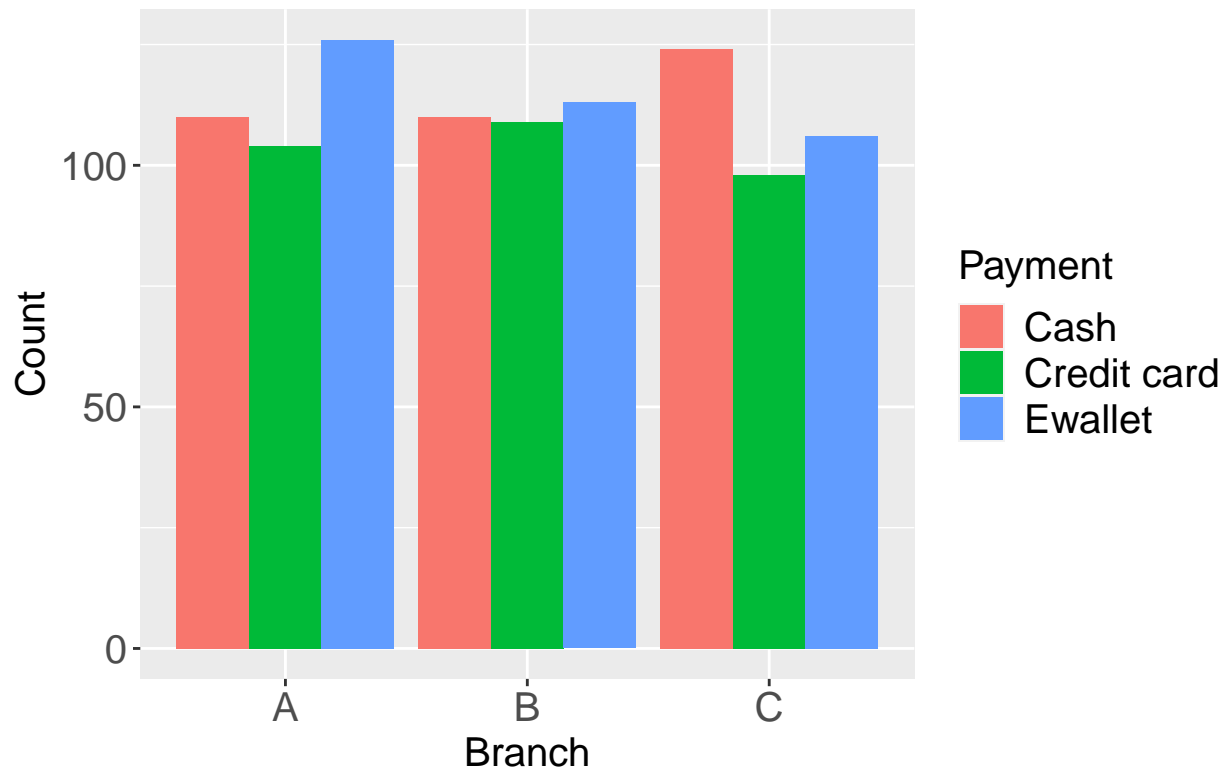
males were known to purchase electronic accessories, health and beauty products and home and lifestyle products females were known to purchase more of fashion and accessories, food and beverages, sports and travel

```
# Plot between payment method and branch

options(repr.plot.width = 15, repr.plot.height = 8)
ggplot(df, aes(x = Branch, fill = Payment)) +
  geom_bar(position = "dodge") +
  labs(title = "Branch vs payment method", x = "Branch", y = "Count") +
  theme(axis.text = element_text(size=15),
        axis.title = element_text(size = 15),
        plot.title = element_text(hjust = 0.5, size = 22),
        legend.title = element_text(size=15),

        legend.text = element_text(size=15))
```

Branch vs payment method



Cash was most used in branch C, credit card in branch B and e-wallets in branch A

```
# Loading package ggcorrplot
```

```
library(ggcorrplot)
```

```
# Finding the covariance
```

```
cov(numeric)
```

```
##          Unit.price    Quantity      Tax      cogs
## Unit.price    701.9653313  0.83477848  196.6683401  3933.36680
## Quantity      0.8347785   8.54644645   24.1495704   482.99141
## Tax           196.6683401  24.14957038  137.0965941  2741.93188
## cogs          3933.3668019  482.99140761  2741.9318829  54838.63766
## gross.margin.percentage  0.0000000  0.00000000  0.0000000  0.00000
## gross.income    196.6683401  24.14957038  137.0965941  2741.93188
## Rating          -0.3996675  -0.07945646  -0.7333003  -14.66601
## Total          4130.0351420  507.14097799  2879.0284770  57580.56954
## gross.margin.percentage gross.income      Rating
## Unit.price              0  196.6683401  -0.39966752
## Quantity                0   24.1495704  -0.07945646
## Tax                    0  137.0965941  -0.73330028
## cogs                   0  2741.9318829 -14.66600553
## gross.margin.percentage  0   0.0000000  0.00000000
## gross.income            0  137.0965941  -0.73330028
## Rating                  0  -0.7333003   2.95351823
## Total                  0  2879.0284770 -15.39930581
##          Total
```



```
## Unit.price          4130.03514
## Quantity            507.14098
## Tax                 2879.02848
## cogs                57580.56954
## gross.margin.percentage 0.00000
## gross.income        2879.02848
## Rating              -15.39931
## Total               60459.59802
```

```
# Finding the correlation matrix
cor(numeric)
```


```
## Warning in cor(numeric): the standard deviation is zero
```

```
##           Unit.price  Quantity    Tax      cogs
## Unit.price  1.000000000  0.01077756  0.6339621  0.6339621
## Quantity    0.010777564  1.00000000  0.7055102  0.7055102
## Tax         0.633962089  0.70551019  1.0000000  1.0000000
## cogs        0.633962089  0.70551019  1.0000000  1.0000000
## gross.margin.percentage NA      NA      NA      NA
## gross.income 0.633962089  0.70551019  1.0000000  1.0000000
## Rating      -0.008777507 -0.01581490 -0.0364417 -0.0364417
## Total       0.633962089  0.70551019  1.0000000  1.0000000
##           gross.margin.percentage gross.income    Rating
## Unit.price                NA      0.6339621 -0.008777507
## Quantity                  NA      0.7055102 -0.015814905
## Tax                       NA      1.0000000 -0.036441705
## cogs                      NA      1.0000000 -0.036441705
## gross.margin.percentage    1      NA      NA
## gross.income              NA      1.0000000 -0.036441705
## Rating                    NA     -0.0364417  1.000000000
## Total                     NA      1.0000000 -0.036441705
##           Total
## Unit.price  0.6339621
## Quantity    0.7055102
## Tax         1.0000000
## cogs        1.0000000
## gross.margin.percentage NA
## gross.income 1.0000000
## Rating      -0.0364417
## Total       1.0000000
```

```
# Plotting the correlation matrix
options(repr.plot.width = 15, repr.plot.height = 10)
ggcorrplot(cor(numeric), tl.cex = 15) +
  labs(title = "Correlation Heatmap") +
  theme(axis.title = element_text(size = 12),
        plot.title = element_text(hjust = 0.5, size = 22),
        legend.title = element_text(size=12),
        legend.text = element_text(size=12))
```

```
## Warning in cor(numeric): the standard deviation is zero
```

Correlation Heatmap



Feature Selection

Filter method

```
# Getting the structure of the columns
str(df1)
```

```
## 'data.frame': 1000 obs. of 13 variables:
## $ Branch : chr "A" "C" "A" "A" ...
## $ Customer.type : chr "Member" "Normal" "Normal" "Member" ...
## $ Gender : chr "Female" "Female" "Male" "Male" ...
## $ Product.line : chr "Health and beauty" "Electronic accessories" "Home and lifestyle" ...
## $ Unit.price : num 74.7 15.3 46.3 58.2 86.3 ...
## $ Quantity : int 7 5 7 8 7 7 6 10 2 3 ...
## $ Tax : num 26.14 3.82 16.22 23.29 30.21 ...
## $ Payment : chr "Ewallet" "Cash" "Credit card" "Ewallet" ...
## $ cogs : num 522.8 76.4 324.3 465.8 604.2 ...
## $ gross.margin.percentage: num 4.76 4.76 4.76 4.76 4.76 ...
## $ gross.income : num 26.14 3.82 16.22 23.29 30.21 ...
## $ Rating : num 9.1 9.6 7.4 8.4 5.3 4.1 5.8 8 7.2 5.9 ...
## $ Total : num 549 80.2 340.5 489 634.4 ...
```

```
# Encoding the branch, customer type, Gender, Payment and quantity columns
```

```
df1$Branch = ifelse(df1$Branch == "A", 0,
                    ifelse(df1$Branch == "B", 1,2))
```

```

df1$Customer.type = ifelse(df1$Customer.type == "Normal", 0, 1)
df1$Gender = ifelse(df1$Gender == "Male", 0, 1)
df1$Payment = ifelse(df1$Payment == "Cash", 0,
                     ifelse(df1$Payment == "Credit card", 1,2))

df1$Product.line = ifelse(df1$Product.line == "Electronic accessories", 0,
                          ifelse(df1$Product.line == "Fashion accessories", 1,
                          ifelse(df1$Product.line == "Food and beverages", 2,
                          ifelse(df1$Product.line == "Health and beauty", 3,
                          ifelse(df1$Product.line == "Home and lifestyle", 4, 5))))))

head(df1)

```

```

##      Branch Customer.type Gender Product.line Unit.price Quantity      Tax Payment
## 1         0             1      1           3       74.69         7 26.1415      2
## 2         2             0      1           0       15.28         5  3.8200      0
## 3         0             0      0           4       46.33         7 16.2155      1
## 4         0             1      0           3       58.22         8 23.2880      2
## 5         0             0      0           5       86.31         7 30.2085      2
## 6         2             0      0           0       85.39         7 29.8865      2
##      cogs gross.margin.percentage gross.income Rating      Total
## 1 522.83              4.761905      26.1415     9.1 548.9715
## 2  76.40              4.761905       3.8200     9.6  80.2200
## 3 324.31              4.761905      16.2155     7.4 340.5255
## 4 465.76              4.761905      23.2880     8.4 489.0480
## 5 604.17              4.761905      30.2085     5.3 634.3785
## 6 597.73              4.761905      29.8865     4.1 627.6165

```

Loading the caret and mlbench packages

```

library(mlbench)
library(caret)

```

Loading required package: lattice

Filtering the numerical columns

```

num = df1[, c(5,7,9,11,12)]
head(num)

```

```

##      Unit.price      Tax   cogs gross.income Rating
## 1       74.69 26.1415 522.83      26.1415     9.1
## 2       15.28  3.8200  76.40       3.8200     9.6
## 3       46.33 16.2155 324.31      16.2155     7.4
## 4       58.22 23.2880 465.76      23.2880     8.4
## 5       86.31 30.2085 604.17      30.2085     5.3
## 6       85.39 29.8865 597.73      29.8865     4.1

```

Determining the correlated features

```

set.seed(5)

```

```

# calculate correlation matrix
correlationMatrix <- cor(num)

```

find attributes that are highly corrected (ideally >0.75)

```
highlyCorrelated <- findCorrelation(correlationMatrix, cutoff=0.75)
```

```
# print indexes of highly correlated attributes  
names(num[, highlyCorrelated])
```

```
## [1] "Tax" "cogs"
```

The Tax, cogs are considered the redundant columns since they are highly correlated

```
# Removing the redundant columns and gross percentage column
```

```
df2 <- num[-highlyCorrelated]  
head(df2)
```

```
##   Unit.price gross.income Rating  
## 1      74.69      26.1415    9.1  
## 2      15.28       3.8200    9.6  
## 3      46.33      16.2155    7.4  
## 4      58.22      23.2880    8.4  
## 5      86.31      30.2085    5.3  
## 6      85.39      29.8865    4.1
```


Feature Selection

```
# Performing our graphical comparison
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
par(mfrow = c(1, 2))  
corrplot(correlationMatrix, order = "hclust")  
corrplot(cor(num), order = "hclust")
```



```
# Checking if there is a zero column to unit variance
# Dropping gross.margin.percentage column since it has a standard deviation of 0
```

```
df3 <- within(num, rm(gross.margin.percentage))
```

```
## Warning in rm(gross.margin.percentage): object 'gross.margin.percentage' not
## found
```

```
head(df2)
```

```
##   Unit.price gross.income Rating
## 1    74.69    26.1415    9.1
## 2    15.28     3.8200    9.6
## 3    46.33    16.2155    7.4
## 4    58.22    23.2880    8.4
## 5    86.31    30.2085    5.3
## 6    85.39    29.8865    4.1
```

```
# Pass df to the prcomp(). We also set two arguments, center and scale
```

```
df.pca <- prcomp(df3, center = TRUE, scale. = TRUE)
summary(df.pca)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5
## Standard deviation  1.8673 0.9996 0.7171 2.415e-16 1.533e-16
## Proportion of Variance 0.6973 0.1998 0.1028 0.000e+00 0.000e+00
## Cumulative Proportion 0.6973 0.8972 1.0000 1.000e+00 1.000e+00
```

As a result we obtain 5 principal components, each which explain a percentage of the total variation of the dataset. PC1 explains 69.7% of the total variance, which means that nearly two-thirds of the information in the dataset (5 variables) can be encapsulated by just that one Principal Component. PC2 and 3 explains 29% of the variance.

```
# Calling str() to have a look at your PCA object
```

```
# ---
```

```
#
```

```
str(df.pca)
```

```
## List of 5
```

```
## $ sdev      : num [1:5] 1.87 1.00 7.17e-01 2.41e-16 1.53e-16
```

```
## $ rotation: num [1:5, 1:5] -0.4039 -0.528 -0.528 -0.528 0.0246 ...
```

```
## ..- attr(*, "dimnames")=List of 2
```

```
## .. ..$ : chr [1:5] "Unit.price" "Tax" "cogs" "gross.income" ...
```

```
## .. ..$ : chr [1:5] "PC1" "PC2" "PC3" "PC4" ...
```

```
## $ center   : Named num [1:5] 55.67 15.38 307.59 15.38 6.97
```

```
## ..- attr(*, "names")= chr [1:5] "Unit.price" "Tax" "cogs" "gross.income" ...
```

```
## $ scale    : Named num [1:5] 26.49 11.71 234.18 11.71 1.72
```

```
## ..- attr(*, "names")= chr [1:5] "Unit.price" "Tax" "cogs" "gross.income" ...
```

```
## $ x        : num [1:1000, 1:5] -1.7153 2.2171 0.0354 -1.0882 -2.4971 ...
```

```
## ..- attr(*, "dimnames")=List of 2
```

```
## .. ..$ : NULL
```

```
## .. ..$ : chr [1:5] "PC1" "PC2" "PC3" "PC4" ...
```

```
## - attr(*, "class")= chr "prcomp"
```

```
# Plotting our pca. This will provide us with some very useful insights i.e.
```

```
# Installing our ggbiplot visualisation package
```

```
library(devtools)
```

```
## Loading required package: usethis
```

```
install_github("vqv/ggbiplot")
```

```
## WARNING: Rtools is required to build R packages, but is not currently installed.
```

```
##
```

```
## Please download and install Rtools 4.2 from https://cran.r-project.org/bin/windows/Rtools/ or https://cran.r-project.org/bin/windows/Rtools/
```

```
## Skipping install of 'ggbiplot' from a github remote, the SHA1 (7325e880) has not changed since last
```

```
## Use `force = TRUE` to force installation
```

```
# Loading our ggbiplot library
```

```
library(ggbiplot)
```

```
## Loading required package: plyr
```

```
## -----
```

```
## You have loaded plyr after dplyr - this is likely to cause problems.
```

```
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
```

```
## library(plyr); library(dplyr)
```

```
## -----
```

```
##
```

```
## Attaching package: 'plyr'
```


```
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
```

```
## Loading required package: scales
```

```
## Loading required package: grid
```

```
# Visualizing our plot
```

```
ggbiplot(df.pca)
```



```
# Adding more detail to the plot, we provide arguments row names as labels
```

```
ggbiplot(df.pca, labels=rownames(df.pca), obs.scale = 1, var.scale = 1)
```

